

# Waze Project

Milestone 5 / 5a - Regression analysis: Build a regression model. Evaluate the model

## Regression modeling

**The purpose** of this project is to conduct exploratory data analysis (EDA) and build a binomial logistic regression model.

**The goal** is to build a binomial logistic regression model and evaluate the model's performance.

*This notebook has three parts:*

**Part 1:** EDA & Checking Model Assumptions

**Part 2:** Model Building, Results, and Evaluation

**Part 3:** Conclusions, Insights, and Recommendations

## Imports and data loading

```
In [1]: # Packages for numerics + dataframes
import pandas as pd
import numpy as np

# Packages for visualization
import matplotlib.pyplot as plt
import seaborn as sns

# Packages for Logistic Regression & Confusion Matrix
from sklearn.preprocessing import StandardScaler, OneHotEncoder
from sklearn.model_selection import train_test_split
from sklearn.metrics import classification_report, accuracy_score, precision_score, \
recall_score, f1_score, confusion_matrix, ConfusionMatrixDisplay
from sklearn.linear_model import LogisticRegression
```

```
In [4]: # Load the dataset by running this cell
df = pd.read_csv('https://raw.githubusercontent.com/adacert/waze/main/Synthetic_Waze_Dat
```

## Part 1. Explore data with EDA & Checking model assumptions

```
In [5]: print(df.shape)

df.info()

(14999, 13)
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 14999 entries, 0 to 14998
Data columns (total 13 columns):
 #   Column              Non-Null Count  Dtype  
---  -
 0   ID                  14999 non-null  int64
```

```

1   label      14299 non-null object
2   sessions   14999 non-null int64
3   drives     14999 non-null int64
4   total_sessions 14999 non-null float64
5   n_days_after_onboarding 14999 non-null int64
6   total_navigations_fav1 14999 non-null int64
7   total_navigations_fav2 14999 non-null int64
8   driven_km_drives 14999 non-null float64
9   duration_minutes_drives 14999 non-null float64
10  activity_days 14999 non-null int64
11  driving_days 14999 non-null int64
12  device      14999 non-null object
dtypes: float64(3), int64(8), object(2)
memory usage: 1.5+ MB

```

The label column is missing 700 values

```
In [6]: df.head()
```

```

Out[6]:
   ID  label  sessions  drives  total_sessions  n_days_after_onboarding  total_navigations_fav1  total_navigati
0  0  retained      283    226      296.748273                2276                208
1  1  retained      133    107      326.896596                1225                 19
2  2  retained      114     95      135.522926                2651                 0
3  3  retained       49     40       67.589221                 15                322
4  4  retained       84     68      168.247020                1562                166

```

Remove the ID column since we don't need this information.

```
In [7]: df = df.drop('ID', axis=1)
```

Class balance of the dependent (target) variable, `label` .

```
In [8]: df['label'].value_counts(normalize=True)
```

```

Out[8]:
retained    0.822645
churned     0.177355
Name: label, dtype: float64

```

```
In [9]: df.describe()
```

```

Out[9]:
      sessions      drives  total_sessions  n_days_after_onboarding  total_navigations_fav1  total_naviga
count  14999.000000  14999.000000   14999.000000             14999.000000             14999.000000             14
mean    80.633776    67.281152    189.964447             1749.837789             121.605974
std    80.699065    65.913872    136.405128             1008.513876             148.121544
min      0.000000      0.000000      0.220211              4.000000              0.000000
25%     23.000000    20.000000     90.661156             878.000000              9.000000
50%     56.000000    48.000000    159.568115             1741.000000             71.000000
75%    112.000000    93.000000    254.192341             2623.500000             178.000000
max    743.000000   596.000000   1216.154633             3500.000000            1236.000000

```

The following columns all seem to have outliers:

sessions, drives, total\_sessions, total\_navigations\_fav1, total\_navigations\_fav2, driven\_km\_drives, duration\_minutes\_drives

The maximum values of all these columns surpass the 75th percentile by multiple standard deviations, suggesting the presence of potential outliers in these variables.

## Create features

```
In [10]: # 1. Create `km_per_driving_day` column
df['km_per_driving_day'] = df['driven_km_drives'] / df['driving_days']

# 2. Call `describe()` on the new column
df['km_per_driving_day'].describe()
```

```
Out[10]: count      1.499900e+04
mean              inf
std              NaN
min       3.022063e+00
25%       1.672804e+02
50%       3.231459e+02
75%       7.579257e+02
max              inf
Name: km_per_driving_day, dtype: float64
```

Note that some values are infinite. This is the result of there being values of zero in the `driving_days` column.

```
In [11]: # 1. Convert infinite values to zero
df.loc[df['km_per_driving_day']==np.inf, 'km_per_driving_day'] = 0

# 2. Confirm that it worked
df['km_per_driving_day'].describe()
```

```
Out[11]: count      14999.000000
mean         578.963113
std         1030.094384
min           0.000000
25%         136.238895
50%         272.889272
75%         558.686918
max        15420.234110
Name: km_per_driving_day, dtype: float64
```

### professional\_driver

Creates a new, binary feature called `professional_driver` that is a 1 for users who had 100 or more drives and drove on 20+ days in the last month.

**Note:** The objective is to create a new feature that separates professional drivers from other drivers.

```
In [12]: # Create `professional_driver` column
df['professional_driver'] = np.where((df['drives'] >= 60) & (df['driving_days'] >= 15),
```

```
In [13]: # 1. Check count of professionals and non-professionals
print(df['professional_driver'].value_counts())

# 2. Check in-class churn rate
df.groupby(['professional_driver'])['label'].value_counts(normalize=True)
```

```
0      12405
```



<b>25%</b>	23.000000	20.000000	90.457733	878.500000	10.000000
<b>50%</b>	56.000000	48.000000	158.718571	1749.000000	71.000000
<b>75%</b>	111.000000	93.000000	253.540450	2627.500000	178.000000
<b>max</b>	243.000000	200.000000	455.439492	3500.000000	422.000000

## Encode categorical variables

```
In [18]: # Create binary `label2` column
df['label2'] = np.where(df['label']=='churned', 1, 0)
df[['label', 'label2']].tail()
```

```
Out[18]:
```

	label	label2
<b>14994</b>	retained	0
<b>14995</b>	retained	0
<b>14996</b>	retained	0
<b>14997</b>	churned	1
<b>14998</b>	retained	0

## Checking assumptions

The following are the assumptions for this logistic regression:

- Independent observations
- No extreme outliers
- Little to no multicollinearity among X predictors
- Linear relationship between X and the **logit** of y

## Collinearity

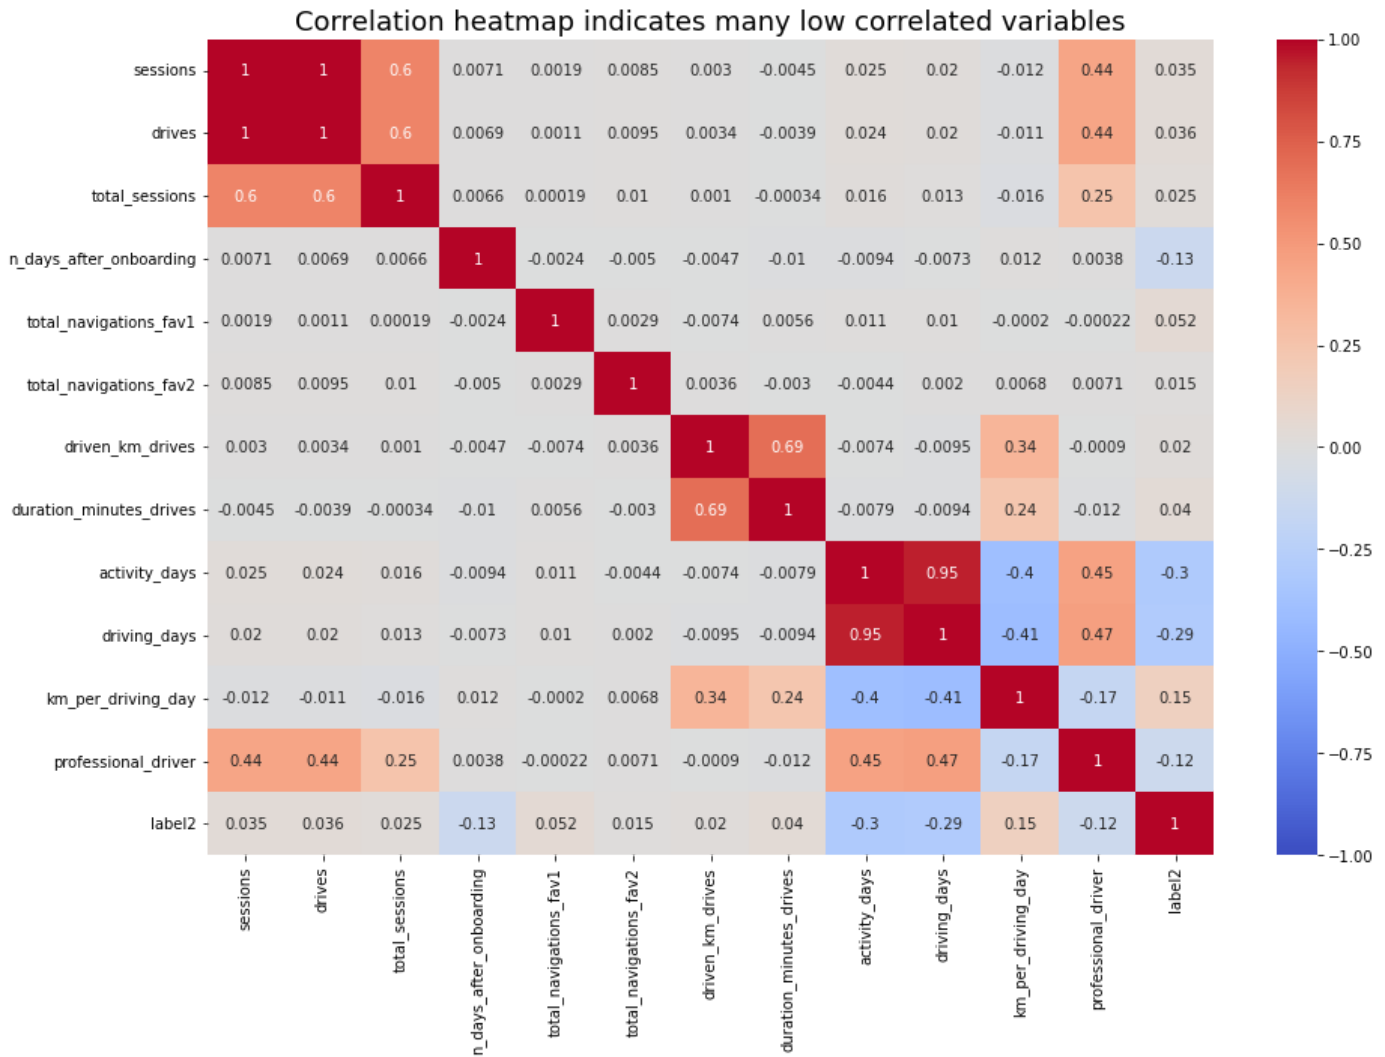
```
In [20]: # Generate a correlation matrix
df.corr(method='pearson')
```

```
Out[20]:
```

	sessions	drives	total_sessions	n_days_after_onboarding	total_navigations_fav1
<b>sessions</b>	1.000000	0.996942	0.597189	0.007101	0.001858
<b>drives</b>	0.996942	1.000000	0.595285	0.006940	0.001058
<b>total_sessions</b>	0.597189	0.595285	1.000000	0.006596	0.000187
<b>n_days_after_onboarding</b>	0.007101	0.006940	0.006596	1.000000	-0.002450
<b>total_navigations_fav1</b>	0.001858	0.001058	0.000187	-0.002450	1.000000
<b>total_navigations_fav2</b>	0.008536	0.009505	0.010371	-0.004968	0.002866
<b>driven_km_drives</b>	0.002996	0.003445	0.001016	-0.004652	-0.007368
<b>duration_minutes_drives</b>	-0.004545	-0.003889	-0.000338	-0.010167	0.005646
<b>activity_days</b>	0.025113	0.024357	0.015755	-0.009418	0.010902
<b>driving_days</b>	0.020294	0.019608	0.012953	-0.007321	0.010419
<b>km_per_driving_day</b>	-0.011569	-0.010989	-0.016167	0.011764	-0.000197

professional_driver	0.443654	0.444425	0.254433	0.003770	-0.000224
label2	0.034911	0.035865	0.024568	-0.129263	0.052322

```
In [22]: # Plot correlation heatmap
plt.figure(figsize=(15,10))
sns.heatmap(df.corr(method='pearson'), vmin=-1, vmax=1, annot=True, cmap='coolwarm')
plt.title('Correlation heatmap indicates many low correlated variables',
          fontsize=18)
plt.show();
```



Variables that are multicollinear with each other?

- sessions and drives: 1.0
- driving\_days and activity\_days: 0.95

## Create dummies

Creates a new, binary column called `device2` that encodes user devices as follows:

- Android -> 0
- iPhone -> 1

```
In [23]: # Create new `device2` variable
df['device2'] = np.where(df['device']=='Android', 0, 1)
df[['device', 'device2']].tail()
```

Out[23]:

	device	device2
14994	iPhone	1
14995	Android	0
14996	iPhone	1
14997	iPhone	1
14998	iPhone	1

## Part 2. Model building, Results, and Evaluation

### Assign predictor variables and target

```
In [24]: # Isolate predictor variables
X = df.drop(columns = ['label1', 'label2', 'device', 'sessions', 'driving_days'])
```

```
In [25]: # Isolate target variable
y = df['label2']
```

### Split the data

```
In [26]: # Perform the train-test split
X_train, X_test, y_train, y_test = train_test_split(X, y, stratify=y, random_state=42)
```

```
In [27]: # Use .head()
X_train.head()
```

Out[27]:

	drives	total_sessions	n_days_after_onboarding	total_navigations_fav1	total_navigations_fav2	driven_kn
152	108	186.192746	3116	243	124	8896
11899	2	3.487590	794	114	18	3286
10937	139	347.106403	331	4	7	7406
669	108	455.439492	2320	11	4	6566
8406	10	89.475821	2478	135	0	1276

### Instantiate a logistic regression model

Add the argument `penalty = None`.

We add `penalty = None` since the predictors are unscaled.

```
In [30]: model = LogisticRegression(penalty='none', max_iter=400)

model.fit(X_train, y_train)
```

```
Out[30]: LogisticRegression(C=1.0, class_weight=None, dual=False, fit_intercept=True,
                             intercept_scaling=1, l1_ratio=None, max_iter=400,
                             multi_class='auto', n_jobs=None, penalty='none',
                             random_state=None, solver='lbfgs', tol=0.0001, verbose=0,
                             warm_start=False)
```

```
In [31]: pd.Series(model.coef_[0], index=X.columns)
```

```
Out[31]: drives          0.001913
total_sessions 0.000327
```

```

n_days_after_onboarding      -0.000406
total_navigations_fav1       0.001232
total_navigations_fav2       0.000931
driven_km_drives              -0.000015
duration_minutes_drives       0.000109
activity_days                 -0.106032
km_per_driving_day            0.000018
professional_driver           -0.001529
device2                       -0.001041
dtype: float64

```

```
In [32]: model.intercept_
```

```
Out[32]: array([-0.00170675])
```

## Check final assumption

Verifies the linear relationship between X and the estimated log odds (known as logits) by making a regplot.

```
In [33]: # Get the predicted probabilities of the training data
training_probabilities = model.predict_proba(X_train)
training_probabilities
```

```
Out[33]: array([[0.93963483, 0.06036517],
               [0.61967304, 0.38032696],
               [0.76463181, 0.23536819],
               ...,
               [0.91909641, 0.08090359],
               [0.85092112, 0.14907888],
               [0.93516293, 0.06483707]])
```

Below creates a dataframe called `logit_data` that is a copy of `df`.

Below also creates a new column called `logit` in the `logit_data` dataframe. The data in this column should represent the logit for each user.

```
In [34]: # 1. Copy the `X_train` dataframe and assign to `logit_data`
logit_data = X_train.copy()


# 2. Create a new `logit` column in the `logit_data` df
logit_data['logit'] = [np.log(prob[1] / prob[0]) for prob in training_probabilities]
```

Below creates a dataframe called `logit_data` that is a copy of `df`.

Below also creates a new column called `logit` in the `logit_data` dataframe. The data in this column should represent the logit for each user.

```
In [35]: # Plot regplot of `activity_days` log-odds
sns.regplot(x='activity_days', y='logit', data=logit_data, scatter_kws={'s': 2, 'alpha':
plt.title('Log-odds: activity_days');
```





## Results and evaluation

If the logistic assumptions are met, the model results can be appropriately interpreted.

Below we will make predictions on the test data.

```
In [36]: # Generate predictions on X_test
y_preds = model.predict(X_test)
```

### Accuracy of the model

```
In [37]: # Score the model (accuracy) on the test data
model.score(X_test, y_test)
```

```
Out[37]: 0.8237762237762237
```

## Results shown with a confusion matrix

```
In [53]: cm = confusion_matrix(y_test, y_preds)
```

The below confusion matrix shows an error, but displays correctly.

```
In [54]: disp = ConfusionMatrixDisplay(confusion_matrix=cm, display_labels=None)
disp.plot()
```

```
-----
TypeError                                Traceback (most recent call last)
<ipython-input-54-5be7a6a26f01> in <module>
      1 disp = ConfusionMatrixDisplay(confusion_matrix=cm, display_labels=None)
----> 2 disp.plot()

/opt/conda/lib/python3.7/site-packages/sklearn/metrics/_plot/confusion_matrix.py in plot
(self, include_values, cmap, xticks_rotation, values_format, ax)
    107         yticklabels=self.display_labels,
    108         ylabel="True label",
--> 109         xlabel="Predicted label")
    110
    111         ax.set_ylim((n_classes - 0.5, -0.5))

/opt/conda/lib/python3.7/site-packages/matplotlib/artist.py in set(self, **kwargs)
    1099         sorted(kwargs.items(), reverse=True,
    1100              key=lambda x: (self._prop_order.get(x[0], 0), x[0])))
```

```

-> 1101         return self.update(props)
      1102
      1103     def findobj(self, match=None, include_self=True):

/opt/conda/lib/python3.7/site-packages/matplotlib/artist.py in update(self, props)
      1004
      1005         with cbook._setattr_cm(self, eventson=False):
-> 1006             ret = [_update_property(self, k, v) for k, v in props.items()]
      1007
      1008             if len(ret):

/opt/conda/lib/python3.7/site-packages/matplotlib/artist.py in <listcomp>(.)
      1004
      1005         with cbook._setattr_cm(self, eventson=False):
-> 1006             ret = [_update_property(self, k, v) for k, v in props.items()]
      1007
      1008             if len(ret):


/opt/conda/lib/python3.7/site-packages/matplotlib/artist.py in _update_property(self, k,
v)
      1001                 raise AttributeError('{!r} object has no property {!r}'
      1002                                     .format(type(self).__name__, k))
-> 1003                 return func(v)
      1004
      1005         with cbook._setattr_cm(self, eventson=False):

/opt/conda/lib/python3.7/site-packages/matplotlib/axes/_base.py in set_yticklabels(self,
labels, fontdict, minor, **kwargs)
      3774             kwargs.update(fontdict)
      3775             return self.yaxis.set_ticklabels(labels,
-> 3776                                             minor=minor, **kwargs)
      3777
      3778     def xaxis_date(self, tz=None):

/opt/conda/lib/python3.7/site-packages/matplotlib/axis.py in set_ticklabels(self, tickla
bels, minor, *args, **kwargs)
      1714         "3.1; passing them will raise a TypeError in Matplotlib 3.3.")
      1715         get_labels = []
-> 1716         for t in ticklabels:
      1717             # try calling get_text() to check whether it is Text object
      1718             # if it is Text, get label content

```

TypeError: 'NoneType' object is not iterable



Precision

```

In [55]: # Calculate precision manually
precision = cm[1,1] / (cm[0, 1] + cm[1, 1])
precision

```

Out[55]: 0.5178571428571429

## Recall

```
In [56]: # Calculate recall manually
recall = cm[1,1] / (cm[1, 0] + cm[1, 1])
recall
```

Out[56]: 0.0914826498422713

## Classification Report

```
In [57]: # Create a classification report
target_labels = ['retained', 'churned']
print(classification_report(y_test, y_preds, target_names=target_labels))
```

	precision	recall	f1-score	support
retained	0.83	0.98	0.90	2941
churned	0.52	0.09	0.16	634
accuracy			0.82	3575
macro avg	0.68	0.54	0.53	3575
weighted avg	0.78	0.82	0.77	3575

Although the model demonstrates reasonable precision, its recall is extremely low, indicating a high number of false negative predictions. Consequently, it fails to identify and capture users who are likely to churn.


## Visual representation of the importance of the model's features

```
In [58]: # Create a list of (column_name, coefficient) tuples
feature_importance = list(zip(X_train.columns, model.coef_[0]))

# Sort the list by coefficient value
feature_importance = sorted(feature_importance, key=lambda x: x[1], reverse=True)
feature_importance
```

```
Out[58]: [('drives', 0.001913369447769776),
 ('total_navigations_fav1', 0.001231754741616306),
 ('total_navigations_fav2', 0.0009314786513814626),
 ('total_sessions', 0.00032707088819142904),
 ('duration_minutes_drives', 0.00010909343558951453),
 ('km_per_driving_day', 1.8223094015325207e-05),
 ('driven_km_drives', -1.4860453424647997e-05),
 ('n_days_after_onboarding', -0.00040647763730561445),
 ('device2', -0.0010412175209008018),
 ('professional_driver', -0.0015285041567402024),
 ('activity_days', -0.10603196504385491)]
```

```
In [59]: # Plot the feature importances
import seaborn as sns
sns.barplot(x=[x[1] for x in feature_importance],
            y=[x[0] for x in feature_importance],
            orient='h')
plt.title('Feature importance');
```



## Part 3: Conclusions, Insights, and Recommendations

### Variables that most influenced the model's prediction:

- Among all the features in the model, "activity\_days" emerged as the most significant one, exhibiting a negative correlation with user churn. This finding is not unexpected since "activity\_days" is highly correlated with "driving\_days," which was already identified during the exploratory data analysis (EDA) to have a negative correlation with churn.

### Variables expected to be stronger predictors than they were:

- During the exploratory data analysis (EDA), it was observed that the user churn rate rose in conjunction with increasing values in "km\_per\_driving\_day." The correlation heatmap in this notebook further confirmed this observation, indicating that this variable exhibited the highest positive correlation with churn among all the predictor variables, surpassing others by a significant margin. Surprisingly, in the model, "km\_per\_driving\_day" ranked as the second-least important variable.

### Why might a variable thought to be important not be important in the model?

- In a multiple logistic regression model, the presence of feature interactions can lead to relationships that may appear counterintuitive. This phenomenon represents both a strength and a weakness of predictive models. On one hand, capturing these interactions enhances the predictive capabilities of the model. On the other hand, it complicates the model's interpretability, making it more challenging to explain the underlying relationships.

### Is it recommended that Waze use this model?

- The usefulness of the model depends on its intended purpose. If the model is employed to inform critical business decisions, its performance may not be sufficiently strong, particularly evident from its low recall score. However, if the model is primarily utilized to guide further exploratory efforts and provide insights, it can still offer value in that context.

### Steps that can be taken to improve this model:

- By leveraging domain knowledge, it is possible to engineer new features aimed at improving predictive signal. In the context of this model, one of the engineered features, namely "professional\_driver," emerged as the third-most influential predictor. Additionally, scaling the predictor variables and

reconstructing the model using different combinations of predictors can be beneficial in minimizing noise stemming from unpromising features.

**Additional features that would be needed to help improve the model:**

- It would be beneficial to possess drive-level specifics for individual users, such as drive times and geographic locations. Furthermore, obtaining more detailed information regarding how users engage with the app would likely provide valuable insights. For instance, understanding the frequency at which they report or confirm road hazard alerts. Finally, having knowledge of the monthly count of distinct starting and ending locations inputted by each driver could offer valuable additional information.